



Generative AI - Conversational Agent

Context

A Generative AI based conversational agent designed to facilitate seamless question-and-answer interactions, summarization, and more within a specified knowledge base. This innovative tool operates with intuitive functionality, enabling users to engage in collaborative exploration and retrieval of information in a natural conversational manner.

Features:

- An agnostic solution covering Databricks, Qdrant vector database and cloud managed services
- Can be deployed into any cloud environment using Terraform scripts(bundled with the product)
- Scalable, robust, flexible and secured solution
- Can evolve and industrialize, easily tailored into a bespoke solution with a minimal coding changes
- Domain independent and adaptable for use across various domains

Deliverables include :

- Design of the demonstration solution, including solution architecture and frontend mock-up
- A mvp deployed in Databricks and Azure
- Detailed documentation, including design of the storyline for each use case



Product Evolution



1. Demonstrate POV

Demonstrate whether the current solution is feasible with Client's requirement.

Demonstrate as it is solution:

- Demonstrate how users can submit a question about a set of text documents and receive an AI-generated short answer
- Demonstrate the high level and low-level solution design
- The solution is agnostic and can be deployed in any cloud with a minimal development effort
- Currently the solution is hosted in Databricks and Azure, the same solution also hosted in a single cloud

2. Perform POC

- Conduct **proof of concept** by
- define the scope
- Identify the environment
- Access data sources (this can be anonymous if performed on Vendor's environment)
- Define a timeline and expected outcome
- **Perform code changes based on the defined scope and measure the solution if this is working as expected**

3. Industrialize and Scale-up

- Industrialize the solution by make it client specific
- Build and unit testing of the end-to-end solution (with new changes) in Client's environment
- Maintain security and network standards
- Deployment to production environment
- setting up governance and maintenance processes, adoption & change management strategy
- identify KPI/metrics for adoption and value generated and build dashboard/report, setup change advisory board
- Training and documentation

Description

Outcome

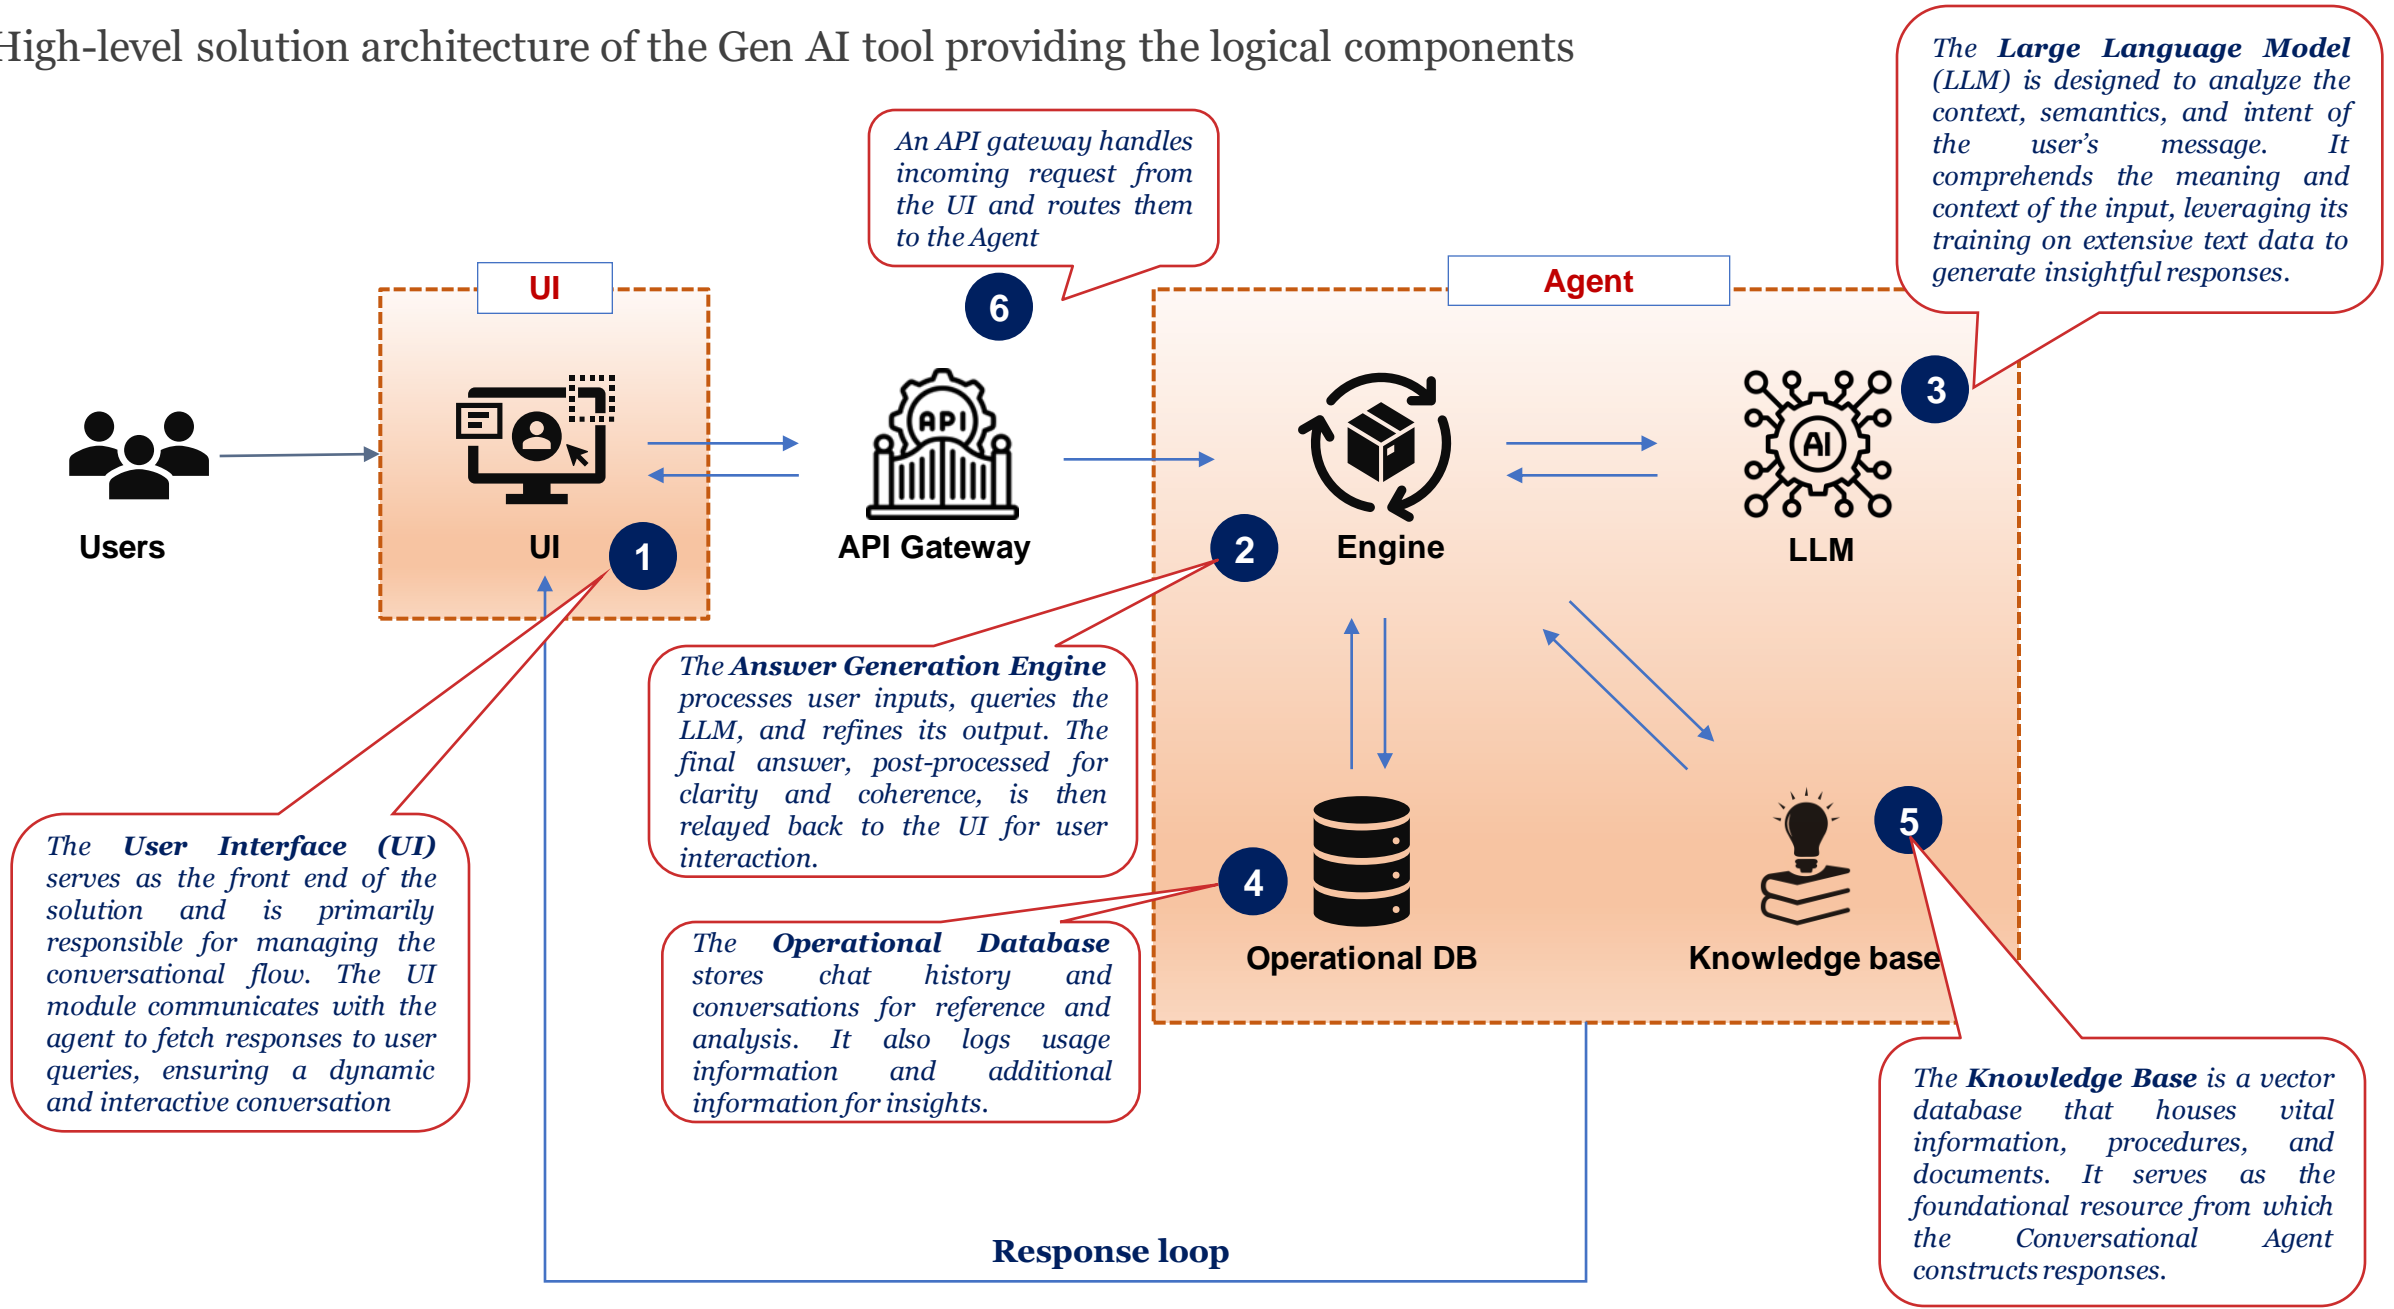
- **Based on Client's interest a POC can be conducted**

- **Based on the POC outcome further evolution can be planned by industrializing the solution**

- **MVP in production**
- **Collect additional requirements/change Request**
- **Continuous evolution and scale-up**

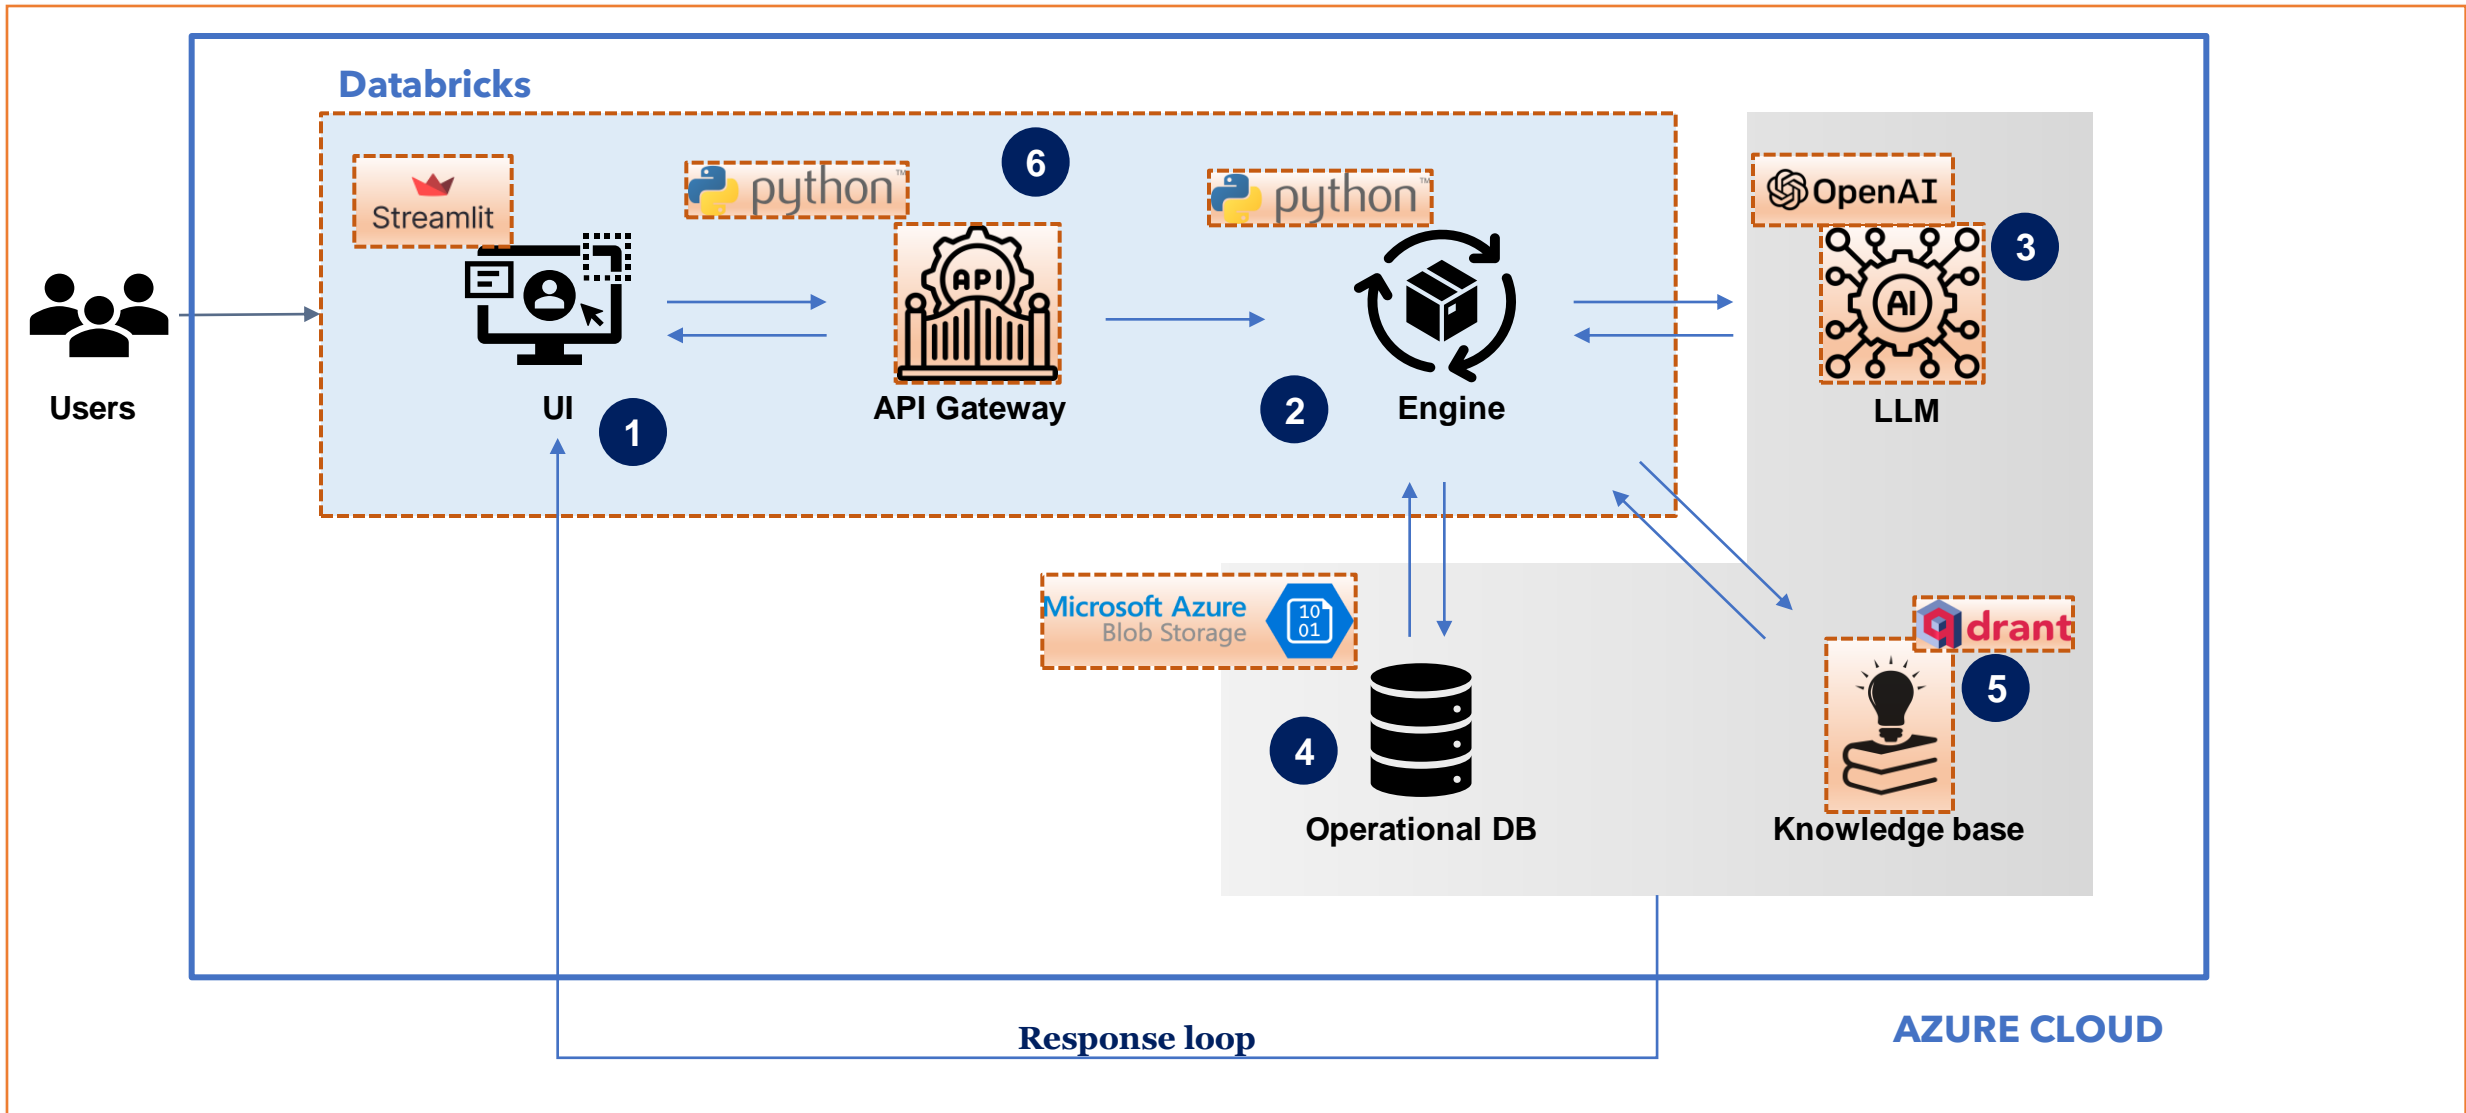
High level solution architecture

High-level solution architecture of the Gen AI tool providing the logical components



Low level solution design (1/2)

A detailed breakdown of the current architecture hosted in Databricks and Azure



Low level solution design

A detailed description of the technology used

Id	Component	Description	Technology	Comments
1	User interface	Manages conversational flow and user interaction	Build a custom user interface hosted in Databricks using Streamlit/react.js	
2	Answer generator engine	Processes user inputs and generates responses	Build a custom Python script hosted in Databricks.	
3	LLM	Analyses user's message context and intent	External AI service managed by the cloud service provider	AWS bedrock/Azure Open AI/ Google Gen AI Studio
4	Operational Database	Stores chat history and logs usage information	No-sql database managed by cloud service provider	Dynamo DB/ Cosmos DB/ Bigtable
5	Knowledge Base	Vector database used to construct responses	Qdrant database hosted on Cloud environment.	A Quadrant database hosted on Kubernetes engine
6	API Gateway	Processes and routes requests to appropriate agents	Build a custom Python script hosted in Databricks notebooks	

A black and white photograph of a hand holding a pen, writing on a document. The background is blurred, showing the document and the hand. The text 'Thank You' is overlaid on the left side of the image.

Thank You

Pragyan Banerjee



LinkedIn



pragyan.banerjee@gmail.com

